

Credit Policy Recommendation Report

To: Chief Risk Officer

From: Engagement Team of Digital Lending Portfolio Optimization

Re: Re-engineering the book toward risk-adjusted growth

Date: 15th June 2026

Classification: **CONFIDENTIAL**

1. Executive Summary

The book is profitable in aggregate (₹47.4 cr on ₹1,398 cr originated, 7.7% 90+ default) but the profit is dangerously concentrated and partly subsidised by a structurally loss-making product. SME is the engine, 20% of loans but ~76% of exposure and ~78% of profit. BNPL is a drag, 35% of loans, near-zero exposure, and loss-making in every product-ticket cell. Personal is the volume backbone, profitable and the product where early-warning intervention pays off most.

We recommend re-engineering the mix rather than tightening underwriting across the board. The headline policy is to contain and re-price BNPL, redirect growth capital into large-ticket SME, and deploy the early-warning system on the Personal book—three moves off existing levers (pricing grid, approval cut-off, channel allocation, collections triage), no new systems required.

Projected net effect: +15–25% book profit and a 0.8-1.5pp reduction in blended 90+ default within 3-4 vintages, at a cost of 8-12% origination volume, almost entirely loss-making BNPL volume we want to give up. Every figure below reconciles to the controlled workbook.

Metric	Value	Metric	Value
Loans	49,600	90+ Default Rate	7.7%
Customers	40,000	Currently Delinquent	15.3%
Originated	₹1,398 cr	WA APR	19.9%
Outstanding	₹651 cr	WA Tenure	23.6 mo
Book Profit	₹47.4 cr	Mean Ticket	₹2.82 lakh

2. Risk Segmentation Findings (Q1)

Four segments separate the book cleanly on risk and value simultaneously. The split is product- and grade-driven, and the value gap between top and bottom is the central finding: the same rupee of capital earns ₹36,974 in SME and *loses* ₹912 in BNPL.

Segment	Defined by	Share of Loans	90+ default	Profit/Loan	Verdict
SME Core (the engine)	SME working capital, grades A–C, larger ticket	~20%	~5.2%	₹36,974 (₹69,604 at >₹1M)	Low risk, very high value → grow
Personal Mainstream	Personal, mid-grades, mid-life seasoning	~45%	~7.0%	₹5,322	Profitable backbone → protect with EWS
BNPL Thin-File (the drag)	BNPL, thin-file, short tenor	~35%	~10.0%	–₹912 (negative at every band)	Structurally loss-making → contain / re-price
Weak-Grade Tail	Grades D–E across products	small	14% (D) → 24% (E)	Negative / marginal	Outsized loss share → price out or cap

The risk gradient is structural, not noise. Default rises monotonically with origination grade (A 1.5% → B 3.5% → C 7% → D 14% → E 24%), so risk is priced-in by grade, not random. Severity runs the opposite way (LGD: SME 75% > Personal 65% > BNPL 40%), so SME's low frequency combined with large ticket makes it the value engine despite the highest loss-given-default with frequency, not severity, is what determines whether a cell makes money here.

Why BNPL cannot be saved by pricing alone. BNPL loses money at every ticket band from <₹10k (–₹917) through >₹25k (–₹803), even at 24-48% APR. The tickets are simply too small to cover acquisition (CAC) and servicing (~₹80/active loan-month). This is a unit-economics problem, not a credit-quality problem, which is why the policy is containment and selective re-pricing rather than tighter scoring.

Value separation is the actionable insight. Plotting the book on a risk $\times$ value matrix yields the classic four quadrants: grow (SME Core), protect (Personal Mainstream), contain/exit (BNPL Thin-File), and re-price/cap (Weak-Grade Tail). Each quadrant maps directly to one policy lever in Section 4.

3. Channel Economics (Q2)

Channel quality differences are mostly a composition effect-aggressive scale channels skew toward marginally weaker borrowers-rather than channels being independently risky. But because BNPL loses money everywhere, the aggregate drag concentrates wherever BNPL is acquired cheapest and at scale. T

Channel	Acquisition Cost (CAC)	Approval TAT	Role
Digital Ads	₹1,222	2-6h	High-volume, scaled- paid and carries the highest BNPL drag (~₹78 Lakhs)
DSA Agent	₹1,021	24-72h	Agent sourced, slower and weaker, second highest drag (~₹29 Lakhs)
Partner-embedded	₹613	1-4h	Embedded checkout, neutral quality
Referral	₹306	6-24h	Referred customers skew better
Organic	₹152	4-24h	Cheapest, self-select, best quality

Implication. The fix is not to shut channels and it is to stop pushing loss-making BNPL through Digital ads and DSA, where the bleed concentrates, and redirect that acquisition budget toward SME and high-quality Personal intake (Referral, Organic, Partner-embedded). This is a reallocation of spend, not a contraction of the funnel.

4. Early-Warning Signal Logic (Q5)

Behavioural deterioration leads to default: for loans that go bad, stress signals begin ramping 2-4 months before the 90+ event. This lead window is what makes early intervention possible and it is strongest on the Personal book, where default is gradual and observable rather than front-loaded as in BNPL.

Target. Transition to 90+ DPD within a 90-day forward window.

Indicators (trend / delta, not level): falling cash-flow consistency, rising balance volatility, first NACH bounce, partial-payment onset, utilisation creep, accelerating bureau-inquiry velocity.

Trigger logic. A governable scorecard (logistic / monotonic gradient-boosted trees) layered with discrete behavioural rules, e.g., two consecutive partial payments + a bounce in the same month banded into a green / amber / red action ladder:

Band	Trigger	Action	Owner
Green	No elevated signal	Monitor only	Automated
Amber	Score elevated / 1 rule fired	Soft contact, restructure offer, watch-list	Collections triage
Red	Score high / multiple rules fired	Priority intervention before roll to 90+	Collections + Risk

Why early action is non-negotiable, the roll-rate evidence. Delinquency cure chance collapses with depth: a loan in the 1-30 bucket has a 68% chance of curing next month, but this falls to 36% at 31-60 and just 21% at 61-90. Critically, 61-90 is the point of no return 74% of loans there roll straight to default, and once a loan reaches 90+ it stays there (93% persistence). The entire value of the EWS is catching loans in the 1-30 / 31-60 window where rescue is still probable. Acting at 61-90 is almost too late.

Expected performance. Median lead time 91 days before default and Mean lead time is 108 days. Thresholds are tied to collections capacity and the cost of a false positive vs. a missed default. The model is strong but deliberately not perfect (AUC = 0.883): 30-40% of stress episodes cure on their own without ever reaching 90+, which is realistic and keeps the false-positive load and collections workload-manageable.

5. Policy Recommendation (Q3, Q4)

Headline policy: *Shift the book from "high-volume, loss-leading BNPL" to "concentrated, high-value SME + EWS-protected Personal," using only existing levers.*

#	Lever	Specific action (who / what / segment / magnitude)	Owner	Guardrail metric
1	Contain & re-price BNPL	Halt origination of loss-making BNPL cells; restrict new BNPL to merchant-subsidised or grade A-B only, or wind down. Cut BNPL spend on Digital ads and DSA first, the two channels carrying ~₹78 lakh and ~₹29 lakh of aggregate drag. Total drag is ~₹1.57 cr.	Product + Credit	BNPL profit / loan $\geq$ ₹0; BNPL share of new loans $\downarrow$
2	Grow large-ticket SME	Redirect freed acquisition capital into SME >₹1M cells (₹69,604 / loan, ~5.2% default). Lift SME origination 3-6% with no loosening of grade cut-offs.	Growth + Credit	SME 90+ default $\leq$ 6%; SME grade-mix stable
3	Risk-based re-pricing + tenure caps on the weak tail	Re-price grades D-E so APR covers expected loss + cost of funds (~11%) + target margin; cap long tenures in weak segments where hazard rises disproportionately with MOB.	Pricing + Credit	Approval-rate floor; volume given up per loss avoided
4	Deploy EWS on Personal	Run the green / amber / red triage live on the Personal book; intervene in the 1-30 / 31-60 window where cure probability is 68% / 36%, with 60-90-day lead.	Risk + Collections	EWS precision vs. collections capacity; share of losses pre-empted

Each lever is operationally feasible inside current systems, has a named owner, and carries a stop/rollback guardrail. Levers are sequenced **1** $\rightarrow$ **4** $\rightarrow$ **3** $\rightarrow$ **2**: contain the bleed and protect the backbone first (fast, low-regret), then re-price and grow (slower, capital-committing).

6. Projected Impact Assessment

Impacts are shown as **ranges with stated assumptions**, not single hero numbers. Baseline = current book (₹47.4 cr profit, 7.7% blended default).

Lever	Δ 90+ default	Δ origination volume	Δ revenue / profit	Net risk-adj. contribution
1. BNPL containment	-0.3 to -0.5pp	-8 to -12% (loss-making)	+₹3-6 cr (loss removed)	+
2. SME large-ticket growth	-0.2 to -0.4pp (mix)	+3 to +6%	+₹5-9 cr	++
3. Re-price weak tail	-0.2 to -0.4pp	-1 to -3%	+₹1-3 cr	+
4. EWS on Personal	losses pre-empted, not rate	neutral	+₹2-4 cr (loss avoided)	+
Combined	-0.8 to -1.5pp	-5 to -9% net	+15 to +25% book profit	strongly +

The growth-vs-risk trade-off is favourable, not painful. The volume given up is overwhelmingly negative-margin BNPL, the firm sheds loans it was paying to make. Net volume falls 5-9%, but profit rises 15–25% and the book gets structurally safer. Sensitivity (tornado, largest first).

BNPL wind-down speed, the single biggest swing. Faster containment lifts profit but costs more headline volume; slower preserves volume optics but prolongs the bleed. SME demand elasticity whether freed capital can actually be deployed into >₹1M SME at the assumed default rate; if SME intake softens, lever 2's upside compresses toward +₹5 cr. EWS precision at the chosen threshold sets how much of the +₹2–4 cr loss-avoidance is realised vs. eaten by false-positive collections cost.

Macro stress (+2–3pp on default). Under stress, the re-priced and EWS-protected book absorbs the shock materially better than baseline; the unmanaged BNPL drag worsens fastest. The recommendation is more valuable in a downturn, not less.

A note on vintages. Read naïvely, recent cohorts look better (raw cumulative default falls toward the newest months) but that is only because recent loans are young and unseasoned. On an equal-age basis (early 30+ by MOB 3), newer cohorts are modestly worse, and grade D+E intake is flat (~24%), so the deterioration is a macro effect, not a riskier grade mix. This reinforces lever 4: the EWS is the right instrument for a book whose recent intake is drifting slightly worse for reasons underwriting alone cannot fix."

Appendix

Data design. Synthetic four-table book mirroring the full customer journey, customers (~40k), loans (~50k), repayments (~381k panel rows), behaviour-monthly (~350k rows); 24 monthly vintages (M1-M24), cut at an M24 snapshot. Seeded and reproducible; passed data-quality, distribution and relationship-conformance gates.

Method. A grade-calibrated monthly hazard with a behavioural stress process that drives *both* default and the early-warning signals (so deterioration is observable before default). Segmentation tested for material separation, interpretability and stability across resamples. All impacts recomputed from the controlled workbook, one number, one source.

Key assumptions. LGD by product, BNPL 40% / Personal 65% / SME 75%. Cost of funds $\approx$ 11% annual. Servicing $\approx$ ₹80 / active loan-month. Fee revenue $\approx$ 1-2% of ticket at origination. Write-off at 180 DPD. Impacts expressed as ranges; macro sensitivity $\pm$ 2-3pp default. EWS target = transition to 90+ within a 90-day forward window; deterioration onset 2-4 months pre-default.

Every figure in this report reconciles to the controlled workbook.